

PEDAGOGICAL REPORT

Advanced Data Visualization Teaching Materials

Course: INFO 7390 – Advanced Data Science and Architecture

Module: Perception-Driven Data Visualization & Storytelling

Author: Mubin Modi

Date: December 2025

EXECUTIVE SUMMARY

This pedagogical report documents the instructional design, theoretical foundations, and implementation decisions for an advanced data visualization teaching module. The module emphasizes evidence-based design principles rooted in cognitive psychology, hands-on practice with industry-standard tools (Plotly, ipywidgets), and authentic assessment through data storytelling. Target audience includes graduate students and professionals in data science with intermediate Python proficiency.

Key Innovations:

- Perception-first instruction grounded in Cleveland & McGill's research
 - Progressive complexity from static charts to interactive storytelling
 - Integration of technical skills with communication abilities
 - Multimodal learning accommodating visual, kinesthetic, and reading/writing styles
-

TABLE OF CONTENTS

1. Teaching Philosophy
 2. Concept Deep Dive
 3. Implementation Analysis
 4. Assessment & Effectiveness
-

1. TEACHING PHILOSOPHY

1.1 Target Audience & Background Assumptions

Primary Audience: Graduate students in data science, computer science, or related fields; professionals transitioning into data science roles; researchers needing to communicate quantitative findings.

Expected Prerequisites:

- Programming: Intermediate Python (functions, loops, data structures)
- Data Manipulation: Basic pandas operations (filtering, grouping, merging)
- Mathematics: High school algebra; understanding of logarithms helpful
- Statistics: Familiarity with mean, median, distribution concepts
- Prior Visualization: Basic matplotlib/seaborn exposure beneficial

Not Required:

- Advanced statistical modeling
- Web development or JavaScript
- Graphic design principles
- Domain expertise in dataset

Accessibility Considerations: Materials designed for diverse learning styles with no assumptions about color vision. Clear, jargon-free explanations with technical terms defined on first use. Scaffolded complexity progressing from beginner to advanced levels.

1.2 Learning Objectives

Knowledge Objectives (Remember/Understand):

1. Explain cognitive science principles underlying effective data visualization (Cleveland & McGill's perceptual hierarchy)
2. Identify when to use logarithmic scales versus linear scales
3. Describe Tufte's data-ink ratio and its importance
4. Recognize common visualization pitfalls (truncated axes, 3D pie charts, misleading scales)

Application Objectives (Apply/Analyze): 5. Select appropriate chart types based on data characteristics and communication goals 6. Construct perception-aware visualizations using Plotly Express 7. Implement interactive widgets for exploratory data analysis 8. Critique existing visualizations using evidence-based design principles

Synthesis/Evaluation Objectives (Create/Evaluate): 9. Design annotated visualizations that tell coherent data stories 10. Evaluate trade-offs between different visualization approaches 11. Create multi-panel narratives with supporting evidence 12. Defend design choices using perceptual and cognitive principles

Bloom's Taxonomy Distribution:

- Create: 35% (Exercise 3: Data Story)
- Evaluate: 25% (Critique sections)
- Analyze: 20% (Implementation choices)
- Apply: 15% (Exercises 1-2)
- Understand: 3% (Conceptual explanations)
- Remember: 2% (Definitions, terminology)

1.3 Pedagogical Approach and Rationale

Core Philosophy: Evidence-Based Design Meets Learning by Doing

This module employs a constructivist learning approach grounded in three principles:

Principle 1: Perception-First Instruction

Rationale: Students often focus on aesthetics ("make it pretty") rather than effectiveness ("make it accurate"). By teaching perceptual science first, students develop a framework for making principled decisions.

Implementation:

- Start with Cleveland & McGill's hierarchy (1984 research)
- Use diagrams showing accuracy differences (bar vs. pie)
- Provide counterexamples of misleading visualizations
- Connect to real-world consequences (misinterpreted COVID-19 charts, financial misinformation)

Why This Works: Grounding design in cognitive science gives students objective criteria rather than subjective taste. Research shows that explicit instruction in perceptual principles improves visualization quality (Hullman et al., 2011).

Principle 2: Scaffolded Code-Along Examples

Rationale: Passive reading doesn't build programming skills. Students need guided practice with immediate feedback.

Implementation Sequence:

- Step 1: Instructor provides complete working code
- Step 2: Students run code, observe output, modify parameters
- Step 3: Students complete partially-filled templates
- Step 4: Students write from scratch with reference materials

Progressive Complexity:

- Static scatter plot (5 parameters)
- Styled scatter plot (10 parameters plus layout)
- Interactive widget (15+ parameters plus callbacks)
- Multi-panel story (20+ parameters plus annotations)

Why This Works: Zone of Proximal Development (Vygotsky) - students work slightly beyond current ability with scaffolding. Gradually removing support builds independence.

Principle 3: Authentic Assessment Through Storytelling

Rationale: Real data scientists must communicate findings to stakeholders. Storytelling integrates technical and communication skills.

Implementation:

- Exercise 3 requires narrative writing plus technical visualization
- Students must identify patterns, not just reproduce examples
- Emphasis on "so what?" rather than "what?"
- Peer review component (if implemented in course)

Why This Works: Transfer of learning improves when assessment mirrors real-world tasks (Wiggins & McTighe, "Understanding by Design"). Storytelling forces synthesis of multiple concepts.

1.4 Alignment with Course Themes

INFO 7390 Course Theme Connections:

Data Architecture: Visualization as final layer in data pipeline (ETL → Analysis → Visualization)

Scalability: Performance considerations for large datasets including sampling, aggregation, and WebGL acceleration

Communication: Translating technical findings for stakeholders through annotation and storytelling frameworks

Best Practices: Evidence-based design versus ad-hoc choices grounded in perceptual research and accessibility standards

Tools & Libraries: Modern Python ecosystem (Plotly, ipywidgets) that is production-ready and interactive

Cross-Cutting Concerns:

- Ethics: Avoiding misleading visualizations (truncated axes, cherry-picking)
- Reproducibility: Jupyter notebooks as executable documentation

- Collaboration: Interactive widgets for exploratory analysis with stakeholders

1.5 Pedagogical Innovations

What Makes This Module Different:

Perception-Driven Rather Than Tool-Driven: Most tutorials teach "how to make X chart." This module teaches "when and why to make X chart." Tool knowledge fades; principles endure.

Interactive Widgets as Pedagogical Tool: Students see cause-effect immediately (slider updates plot). Builds intuition about parameter impacts. Mirrors professional workflows (dashboards, stakeholder demos).

Anti-Patterns as Teaching Moments: Explicitly shows bad visualizations with explanations. Students learn what not to do (often more valuable). Develops critical evaluation skills.

Integrated Storytelling: Not a separate "communication" module. Storytelling woven throughout technical content. Mirrors real-world practice (analysis plus presentation).

2. CONCEPT DEEP DIVE

2.1 Technical & Mathematical Foundations

2.1.1 Perceptual Encoding: Cleveland & McGill's Hierarchy

Historical Context: In 1984, William Cleveland and Robert McGill published groundbreaking research in the Journal of the American Statistical Association establishing an empirical hierarchy of graphical perception.

The Experiment: Subjects estimated proportions from different chart types. Error rates measured for each visual encoding. Results showed position along common scale most accurate (~95%). Worst performers were area, volume, and color saturation (~40-50% accuracy).

Mathematical Foundation - The Weber-Fechner Law: The Weber-Fechner Law from psychophysics explains perceptual accuracy:

$$\Delta I / I = k$$

Where:

- ΔI = just-noticeable difference in stimulus
- I = original stimulus intensity

- k = constant (varies by perceptual channel)

Implication: Human perception of magnitude is ratio-based, not absolute. We're better at comparing positions (linear) than areas (quadratic).

Practical Application: Bar charts use position encoding (most accurate). Pie charts use angle/area encoding (less accurate). For comparing values, horizontal bars superior to pie charts because position along common x-axis enables precise visual comparison.

Code Example:

```
# Bad: Pie chart for comparison (angle/area encoding)
fig = px.pie(df, values='sales', names='product')
```

```
# Good: Horizontal bar for ranking (position encoding)
fig = px.bar(df, x='sales', y='product', orientation='h')
```

2.1.2 Logarithmic Scales: Mathematical Justification

When to Use Log Scales:

A variable X should be log-transformed when:

1. Data spans multiple orders of magnitude: $\max(X) / \min(X) > 100$
2. Interested in multiplicative relationships (percentage growth)
3. Data follows log-normal or power-law distribution

Mathematical Properties:

On a log scale, equal distances represent equal ratios:

$$\log(100) - \log(10) = \log(1000) - \log(100) = 1$$

Both represent $10\times$ growth, despite different absolute changes (90 versus 900).

Visual Impact:

Linear Scale (GDP per capita):

- High values dominate visual space
- Most countries compressed into narrow band
- Difficult to see patterns in lower values

Log Scale:

- All values visible with proportional spacing

- Equal visual space per order of magnitude
- Reveals relationships across full range

Code Implementation:

```
# Plotly automatically handles log transformation
fig = px.scatter(df, x='gdpPercap', y='lifeExp', log_x=True)

# Always label as log scale to prevent misinterpretation
fig.update_xaxes(title="GDP per Capita (log10 scale)")
```

Common Pitfall: Students forget to label as "log scale" causing misinterpretation. Axis labels must explicitly state transformation.

2.1.3 The Data-Ink Ratio: Information Theory Perspective

Tufte's Formulation (1983):

Data-Ink Ratio = (Ink used for data) / (Total ink used)

Goal: Maximize ratio by removing non-data ink (grid lines, borders, decorations).

Information-Theoretic Interpretation:

Chart junk (decorative elements) adds visual complexity without increasing information content. This creates noise in the signal-to-noise ratio:

$SNR = \text{Signal Power} / \text{Noise Power} = \text{Data Information} / \text{Non-Data Distraction}$

Practical Implementation:

```
# Low data-ink ratio (Excel defaults)
fig.update_layout(
    plot_bgcolor='#e0e0e0', # Gray background adds no information
    xaxis=dict(showgrid=True, gridwidth=2), # Heavy gridlines
    yaxis=dict(showgrid=True, gridwidth=2),
    font=dict(size=10) # Small, hard to read
)

# High data-ink ratio (Tufte style)
fig.update_layout(
    plot_bgcolor='white', # Clean background
    xaxis=dict(showgrid=False), # Remove unnecessary grids
    yaxis=dict(showgrid=False),
    font=dict(size=14) # Readable font size
```

)

2.2 Connection to Course Themes

2.2.1 Visualization in the Data Science Pipeline

Traditional Linear View: ETL → Storage → Analysis → Visualization (endpoint)

Actual Iterative Reality: ETL → Storage ⇌ Exploratory Viz → Analysis ⇌ Diagnostic Viz → Communication Viz

Visualization serves three distinct roles:

1. Exploratory (EDA):
 - Purpose: Discover patterns, outliers, distributions
 - Audience: Self (data scientist)
 - Example: Quick scatter matrix, histograms
 - Tool: Interactive widgets, fast iteration
2. Diagnostic (Validation):
 - Purpose: Verify model assumptions, check residuals
 - Audience: Technical team
 - Example: Q-Q plots, residual plots, learning curves
 - Tool: Standardized statistical graphics
3. Communication (Presentation):
 - Purpose: Convince stakeholders, support decisions
 - Audience: Non-technical leaders
 - Example: Annotated narratives, dashboards
 - Tool: Polished, storytelling-focused

Module Focus: Primarily communication (#3) with techniques applicable to exploration (#1) via widgets.

2.2.2 Scalability Considerations

The Big Data Visualization Challenge:

Modern datasets exceed browser rendering limits:

- Plotly.js handles approximately 10K points smoothly
- Greater than 100K points cause lag
- Greater than 1M points crash browser

Solutions Taught:

1. Strategic Sampling:

```
# Random sampling preserves distribution
df_sample = df.sample(n=10000)

# Stratified sampling preserves groups
df_sample = df.groupby('continent').sample(n=2000)
```

2. Aggregation Before Plotting:

```
# Don't plot 1M individual transactions
# Instead aggregate first:
df_agg = df.groupby(['date', 'category']).agg({
    'revenue': 'sum',
    'transactions': 'count'
}).reset_index()
```

3. WebGL Acceleration:

```
# Use Scattergl for greater than 10K points
import plotly.graph_objects as go
fig = go.Figure(data=go.Scattergl(
    x=large_df['x'],
    y=large_df['y'],
    mode='markers'
))
```

Connection to Course: Demonstrates that good visualization requires understanding data volume and computational constraints, not just aesthetics.

2.2.3 Ethical Considerations

How Visualizations Mislead:

1. Truncated Y-Axis (Exaggerates small differences):

```
# Bad: Exaggerates differences
fig.update_yaxes(range=[95, 105])

# Good: Shows true scale
fig.update_yaxes(range=[0, 105])
```

2. Cherry-Picking Time Ranges (Shows only favorable periods):

Bad: Shows only favorable period

```
df_filtered = df[df.year.between(2020, 2022)]
```

Good: Shows full context with annotations for anomalies

Use all available data

3. Dual Y-Axes (Creates spurious correlations):

Can make ANY two variables look correlated by adjusting scales

Example: Ice cream sales and drownings appear correlated

Reality: Both driven by temperature (confounding variable)

Teaching Approach: Show unethical examples explicitly. Discuss why problematic. Provide ethical alternatives. Emphasize transparency (always cite data sources, show full scales).

2.3 Relationship to Real-World Data Science Workflows

2.3.1 Stakeholder Communication Scenario

Realistic Use Case: You're a data scientist at a healthcare company. Your model predicts patient readmission risk. You need to present findings to:

Technical Team: Model architecture, performance metrics

- Precision-recall curves (diagnostic viz)
- Feature importance charts (bar charts with error bars)
- Performance over time (line charts)

Clinical Staff: Which patients to prioritize, why

- Patient risk distribution (histogram with annotations)
- Top risk factors (horizontal bar, sorted)
- Geographic distribution (choropleth map)

Executives: ROI, resource allocation

- Cost savings projection (annotated line chart)
- High-risk patient count (big number plus sparkline)
- Before/after comparison (side-by-side bar chart)

Module Connection: Exercise 3 (storytelling) forces students to create executive-friendly narratives. Widget exercise teaches exploratory tools for technical audiences. Annotation techniques support clinical communication.

2.3.2 Jupyter Notebooks as Executable Reports

Industry Trend: Notebooks increasingly used for:

- Internal reporting: Weekly/monthly data reviews
- Client deliverables: Consulting reports with live data
- Regulatory submissions: Reproducible analysis documentation

Module Implementation:

1. Self-Contained Narratives:

- Markdown cells explain context
- Code cells show methodology
- Output cells present results
- No external dependencies (self-contained PDF export)

2. Reproducibility:

- Version-controlled (Git-friendly)
- Requirements.txt for dependencies
- Data loading from predictable paths
- Random seed setting for consistency

3. Interactivity:

- Widgets allow stakeholders to explore
- Parameter cells at top for easy modification
- Functions abstract complexity

Real-World Pattern:

```
# Executive summary at top
display(Markdown("""
## Q4 2024 Performance Summary
- Revenue: $1.2M (+15% QoQ)
- Customer Acquisition: 450 new users
- Churn Rate: 3.2% (-0.8% QoQ)
"""))

# Interactive deep-dive below
@interact(metric=['revenue', 'users', 'churn'])
def explore(metric):
```

```
fig = px.line(df, x='date', y=metric)
fig.show()
```

3. IMPLEMENTATION ANALYSIS

3.1 Architecture and Design Decisions

3.1.1 Technology Stack Selection

Core Libraries Selected:

Plotly Express 5.x: High-level visualization API

- Rationale: Concise syntax, interactive by default, production-ready

ipywidgets 7.7.1: Interactive controls

- Rationale: Industry standard for Jupyter, rich widget library

Pandas 1.x: Data manipulation

- Rationale: Ubiquitous in Python data science, Plotly integration

NumPy 1.x: Numerical computing

- Rationale: Fast array operations, Plotly backend

Why Plotly Over Alternatives:

Comparison Matrix:

- Interactivity: Plotly (built-in) vs Matplotlib (static, mpld3 clunky) vs Altair (Vega-based) vs D3.js (full control)
- Learning Curve: Plotly (gentle) vs Matplotlib (familiar) vs Altair (grammar-of-graphics) vs D3.js (steep)
- Production Use: Plotly (Dash framework) vs Matplotlib (limited) vs Altair (limited) vs D3.js (custom apps)
- Customization: Plotly (extensive) vs Matplotlib (full control) vs Altair (limited) vs D3.js (unlimited)
- Documentation: Plotly (excellent) vs Matplotlib (mature) vs Altair (good) vs D3.js (scattered)
- Pedagogy: Plotly (quick wins) vs Matplotlib (verbose) vs Altair (declarative) vs D3.js (too complex)

Decision: Plotly Express offers the best effort-to-insight ratio for teaching. Students create publication-quality, interactive visualizations in 5 lines of code.

3.1.2 Module Structure & Information Architecture

Pedagogical Flow:

1. Setup (5%): Imports, data loading, environment check
2. Theory (15%): Perceptual principles, log scales, data-ink ratio
3. Static Examples (30%): Scatter, bar, line charts with explanations
4. Interactive (25%): Widgets, callbacks, dynamic updates
5. Storytelling (15%): Annotations, narratives, story panels
6. Exercises (10%): Progressive difficulty, scaffolded practice

Rationale:

- Front-load theory: Students need mental models before coding
- Immediate application: Each concept followed by code example (theory-practice loop)
- Progressive complexity: Start with 5-line examples, build to 50-line stories
- Capstone integration: Exercise 3 requires synthesizing all prior concepts

Code Organization Pattern:

Concept → Diagram → Code → Output → Discussion

Example Structure:

1. Concept

```
display(Markdown("## Log Scales: When and Why"))
```

2. Diagram (ASCII or image)

```
print("""
```

Linear: Points compressed at low end

Log: Equal visual spacing per order of magnitude

```
""")
```

3. Code (runnable example)

```
fig = px.scatter(df, x='gdpPercap', y='lifeExp', log_x=True)
```

```
fig.show()
```

4. Discussion

```
display(Markdown("""
```

****Why this works:**** GDP spans \$277 to \$113K—four orders of magnitude.

Log scale gives equal visual space to each power of 10.

```
"""))
```

3.1.3 Widget Design Philosophy

Goal: Teach exploratory analysis, not just widget syntax.

Pattern: Observe → Modify → Understand

Step 1: Instructor provides complete interactive example Step 2: Students run, change parameters, observe effects Step 3: Students modify code (add color='continent') Step 4: Students create own widget from scratch

Design Decisions:

1. Widget Type Selection:
 - IntSlider for continuous years (visual feedback)
 - Dropdown for discrete choices (prevents typos)
 - SelectMultiple for multi-select (Ctrl+Click intuitive)
 - Avoid Text input (error-prone, no validation)
2. Callback Pattern:
 - Use @interact decorator (simplest, least code)
 - Alternative: .observe() method (more control, more complex)
 - Teaching progression: decorator → observe → manual updates
3. State Management:
 - Widgets don't persist between notebook sessions
 - Use default values that create meaningful plots
 - Document expected behavior in comments

3.2 Performance Considerations

3.2.1 Rendering Optimization

Challenge: Plotly generates large JSON payloads for complex charts.

Solutions Implemented:

1. Data Downsampling:

```
# For exploration, sample data
if len(df) > 10000:
    df_plot = df.sample(n=10000)
else:
    df_plot = df
```

2. Aggregation:

```
# Don't plot raw transactions; aggregate first
df_agg = df.groupby(['date', 'category']).agg({
    'revenue': 'sum',
    'count': 'size'
}).reset_index()
```

3. Lazy Rendering:

```
# Use Output widget to prevent auto-render
output = widgets.Output()

def update_plot(change):
    with output:
        output.clear_output(wait=True) # Clear before rendering
        fig.show()
```

3.2.2 Memory Management

Issue: Loading large datasets in notebooks can exhaust RAM.

Best Practices Taught:

1. Load Only Needed Columns:

```
# Bad: Loads everything
df = pd.read_csv('huge_file.csv')

# Good: Loads only necessary columns
df = pd.read_csv('huge_file.csv', usecols=['country', 'year', 'lifeExp'])
```

2. Type Optimization:

```
# Convert object to category (saves memory)
df['continent'] = df['continent'].astype('category')

# Downcast numeric types
df['year'] = pd.to_numeric(df['year'], downcast='integer')
```

3. Chunk Processing:

```
# For truly large files
chunks = pd.read_csv('huge_file.csv', chunksize=10000)
df_sample = next(chunks) # Use first chunk for exploration
```

3.3 Edge Cases and Limitations

3.3.1 Data Quality Issues

Problem: Real-world data is messy. Students must handle:

1. Missing Values:

```
# Plotly silently drops NaNs—explain this to students
df_clean = df.dropna(subset=['gdpPercap', 'lifeExp'])
```

```
# Alternative: Fill with median/mean
df['gdpPercap'].fillna(df['gdpPercap'].median(), inplace=True)
```

Teaching Moment: Show what happens when plotting with NaNs (gaps in line charts, missing points in scatter).

2. Outliers:

```
# Extreme values distort scale
# Solution 1: Log scale (taught in module)
# Solution 2: Cap outliers
df['gdpPercap_capped'] = df['gdpPercap'].clip(
    upper=df['gdpPercap'].quantile(0.99)
)
```

```
# Solution 3: Separate visualization for outliers
outliers = df[df['gdpPercap'] > df['gdpPercap'].quantile(0.99)]
```

3. Data Type Mismatches:

```
# Common error: year stored as string
df['year'] = df['year'].astype(int) # Convert for sliders

# Date parsing
df['date'] = pd.to_datetime(df['date_str'], format='%Y-%m-%d')
```

3.3.2 Browser Compatibility

Issue: Interactive notebooks behave differently across environments.

Environment Compatibility Matrix:

- Jupyter Lab: ipywidgets (with extension), Plotly (yes)
- Jupyter Notebook: ipywidgets (built-in), Plotly (yes), Most compatible
- Google Colab: ipywidgets (with setup), Plotly (yes), Requires `output.enable_custom_widget_manager()`
- VSCode: ipywidgets (limited), Plotly (yes), Widgets may not render
- JupyterHub: ipywidgets (admin setup), Plotly (yes), Depends on config

Module Solution:

- Primary platform: Google Colab (zero setup, free GPUs)
- Includes setup code at top of notebook
- Fallback: Static HTML export for viewing without execution

3.3.3 Accessibility Limitations

Known Issues:

1. Screen Readers:
 - Plotly charts not fully accessible to screen readers
 - Mitigation: Provide text summary of key findings

```
display(Markdown("""
```

```
**Chart Summary:** This scatter plot shows 142 countries.
```

```
Top 3 by life expectancy: Japan (82.6), Switzerland (81.7), Iceland (81.8).
```

```
"""))
```

2. Keyboard Navigation:
 - Widgets navigable via Tab, but not obvious
 - Mitigation: Add instructions

```
display(Markdown("Use Tab to navigate widgets, Enter to select"))
```

3. Colorblind-Friendly Palettes:
 - Default Plotly colors not colorblind-safe
 - Solution Taught:

```
# Okabe-Ito colorblind-friendly palette
```

```
colors = ['#E69F00', '#56B4E9', '#009E73', '#F0E442',  
          '#0072B2', '#D55E00', '#CC79A7']
```

```
fig.update_layout(colorway=colors)
```

3.4 Error Handling & Debugging

Common Student Errors:

Data Not Loaded:

```
try:
    df = pd.read_csv('data/gapminder.csv')
except FileNotFoundError:
    print("Warning: Data file not found. Downloading...")
    df = px.data.gapminder() # Fallback to built-in
```

2. Widget Not Displaying:

```
# Check if widgets enabled
try:
    from google.colab import output
    output.enable_custom_widget_manager()
    print("Widgets enabled successfully")
except ImportError:
    print("Warning: Not in Colab. Ensure widgets installed:")
    print("jupyter nbextension enable --py widgetsnbextension")
```

3. Plotly Not Rendering:

```
# Test rendering
fig = px.scatter(x=[1, 2, 3], y=[1, 2, 3])
try:
    fig.show()
    print("Plotly working correctly")
except Exception as e:
    print(f"Plotly error: {e}")
    print("Try: pip install plotly --upgrade")``
```

4. ASSESSMENT & EFFECTIVENESS

4.1 How We Validate Student Understanding

4.1.1 Formative Assessment (During Learning)

Embedded Check-for-Understanding:

1. Conceptual Questions (Throughout Notebook):

Quick Check: Which chart type is best for showing ranking? a) Pie chart b) Scatter plot
c) Horizontal bar chart (CORRECT) d) Line chart

Why? Position along common scale (x-axis) is most accurate perceptual encoding.

2. Code Prediction (Before Running):

Before running this code, predict: Will the points overlap? Will we see all countries?

```
fig = px.scatter(df, x='gdpPercap', y='lifeExp')  
fig.show()
```

Now run it. Were you correct? Why/why not?

3. Error Diagnosis:

This code has a bug. Can you spot it?

```
# x and y are swapped for horizontal orientation  
fig = px.bar(df, x='country', y='lifeExp', orientation='h')  
fig.show()
```

4.1.2 Summative Assessment (Exercises)

Tiered Difficulty Structure:

Exercise 1 (Beginner Level):

- Bloom's Level: Apply
- Skills Assessed: Chart selection, basic Plotly syntax
- Success Criteria: Correct chart type, proper orientation, labeled axes

Exercise 2 (Intermediate Level):

- Bloom's Level: Analyze
- Skills Assessed: Widget implementation, animation trade-offs
- Success Criteria: Working interactive element, written critique of limitations

Exercise 3 (Advanced Level):

- Bloom's Level: Create/Evaluate

- Skills Assessed: Storytelling, annotation, synthesis
- Success Criteria: Coherent narrative, supporting visualizations, design justification

Grading Rubric for Exercise 3 (Storytelling):

Data Insight:

- Excellent (4): Novel, non-obvious pattern identified
- Good (3): Valid pattern, somewhat expected
- Needs Improvement (2): Trivial observation
- Insufficient (1): No clear finding

Visualization Quality:

- Excellent (4): Chart type optimal, annotations clear
- Good (3): Appropriate chart, minor issues
- Needs Improvement (2): Questionable chart choice
- Insufficient (1): Misleading/unclear

Narrative Structure:

- Excellent (4): Clear setup-finding-insight-action
- Good (3): Logical flow, missing one element
- Needs Improvement (2): Disjointed, hard to follow
- Insufficient (1): No coherent story

Design Justification:

- Excellent (4): Cites perceptual principles, defends choices
- Good (3): Some justification provided
- Needs Improvement (2): Weak reasoning
- Insufficient (1): No explanation

Technical Execution:

- Excellent (4): Code runs, no errors, well-commented
- Good (3): Runs with minor warnings
- Needs Improvement (2): Errors present but fixable
- Insufficient (1): Doesn't execute

4.1.3 Self-Assessment Tools

Reflection Prompts:

After each major section:

Reflect:

1. What was the most surprising thing you learned?
2. Which concept do you still find confusing?
3. How would you explain this to a non-technical friend?

Metacognitive Scaffolding:

Before attempting Exercise 3, plan your approach:

1. What question will you investigate?
2. What chart types will you use?
3. What story will you tell?

After completing: 4. What worked well? 5. What would you change?

4.2 Common Challenges Students Face

4.2.1 Conceptual Challenges

Challenge 1: Log Scales Are Unintuitive

Symptom: Students use log scales incorrectly (e.g., when data doesn't span orders of magnitude).

Root Cause: Linear thinking deeply ingrained; multiplicative relationships unfamiliar.

Pedagogical Solution:

1. Anchor to familiar example: "Richter scale (earthquakes) is logarithmic"
2. Visual comparison: Show same data on linear vs. log side-by-side
3. Interactive exploration: Widget to toggle log scale on/off
4. Decision tree: "Use log when max/min > 100"

Interactive Demonstration:

```
@interact(use_log=widgets.Checkbox(description='Use Log Scale'))
def demo_log(use_log):
    fig = px.scatter(df, x='gdpPercap', y='lifeExp', log_x=use_log)
    fig.show()
```

Challenge 2: Choosing Chart Types

Symptom: Students default to scatter plots for everything or use pie charts for rankings.

Root Cause: Limited mental model of chart type to data type mappings.

Pedagogical Solution:

1. Decision tree poster: Flowchart in tutorial documentation
2. Anti-patterns: Explicitly show bad examples (3D pie, dual y-axis)
3. Practice scenarios: "You have categorical data for ranking. What chart?" Answer: "Horizontal bar"

Challenge 3: Interactivity vs. Static

Symptom: Students add widgets unnecessarily (overengineering) or assume all visualizations must be interactive.

Root Cause: Misunderstanding of audience/purpose.

Pedagogical Solution:

1. Context-based teaching: "Exploratory = interactive, Presentation = annotated static"
2. Real-world examples: Show conference slides (static) vs. internal dashboards (interactive)
3. Cost-benefit analysis: "Interactivity takes 3× longer. Worth it for this use case?"

4.2.2 Technical Challenges

Challenge 4: Python Syntax Errors

Symptom: IndentationError, KeyError, AttributeError

Pedagogical Solution:

1. Pre-filled templates: Students fill in blanks rather than write from scratch
2. Incremental building: Start with working code, add one feature at a time
3. Debugging checklist:
 - Did you spell column names exactly? (Case-sensitive)
 - Are all parentheses/brackets matched?
 - Did you run all cells in order?

Challenge 5: Widget State Management

Symptom: "My widget doesn't update the plot" or "Plot renders twice"

Root Cause: Event-driven programming is unfamiliar paradigm.

Pedagogical Solution:

1. Use @interact decorator first (hides complexity)
2. Gradual reveal: Only introduce .observe() pattern for advanced students
3. Visual diagram showing callback flow

Challenge 6: Data Filtering Logic

Symptom: KeyError, empty plots, incorrect subset

Pedagogical Solution:

1. Always print shape after filtering:

```
filtered = df[df.year == 2007]
print(f"Rows after filtering: {len(filtered)}") # Sanity check
```

2. Use .isin() for multiple values:

```
# Common error: df[df.continent == ['Asia', 'Europe']] # Wrong
# Correct:
df[df.continent.isin(['Asia', 'Europe'])] # Right
```

4.3 Addressing Different Learning Styles

4.3.1 Visual Learners (60% of students)


Strengths: Benefit from diagrams, color-coding, spatial organization

Module Accommodations:

- ASCII diagrams in markdown cells
- Color-coded comments (Good vs. Bad examples)
- Side-by-side comparisons (before/after)
- Flowcharts for decision-making (when to use log scale)

Example Visual Aid:

Perception Hierarchy:



4.3.2 Kinesthetic Learners (30% of students)

Strengths: Learn by doing, manipulating, experimenting

Module Accommodations:

- Interactive widgets (immediate feedback)
- "Modify this code and observe" prompts
- Hands-on exercises after each concept
- Encourage parameter experimentation

Example Interactive Exploration:

EXPERIMENT: Change these parameters and observe:

```
fig = px.scatter(  
    df,  
    x='gdpPercap',  
    y='lifeExp',  
    size='pop',      # Try: Remove this line  
    color='continent', # Try: Change to 'year'  
    log_x=True      # Try: Change to False  
)  
fig.show()
```

4.3.3 Reading/Writing Learners (10% of students)

Strengths: Prefer text-based explanations, documentation, written reflection

Module Accommodations:

- Detailed markdown explanations
- Inline code comments
- Reflection prompts (written responses)
- Links to external reading (Tufte, Cleveland papers)

Example Text-Heavy Section:

Deep Dive: Why Log Scales Work

The Weber-Fechner Law from psychophysics explains why humans perceive magnitude ratios more accurately than absolute differences. On a logarithmic scale, equal visual distances represent equal proportional changes...

Further Reading:

- Cleveland & McGill (1984). "Graphical Perception"
- Tufte (1983). "The Visual Display of Quantitative Information"

4.4 Future Improvements and Extensions

4.4.1 Short-Term Enhancements (Next Iteration)

- 1. Video Walkthroughs** Record 5-10 minute concept explainers. Show live coding with narration. Addresses students who missed lecture or need review.
- 2. Auto-Graded Quizzes** Integrate nbgrader for automated exercise checking. Provide instant feedback on correctness. Addresses scalability (class size >100).
- 3. Peer Review Component** Students critique each other's Exercise 3 submissions. Use structured rubric (Section 4.1.2). Addresses communication skills and critical thinking.
- 4. Interactive Debugging Guide** Clickable "I'm getting this error..." troubleshooting. Links to relevant section or documentation. Addresses common technical stumbling blocks.

4.4.2 Medium-Term Extensions (Future Semester)

- 1. Dashboard Deployment Module** Add final project: Deploy Dash app to Heroku. Teaches production skills and stakeholder demos. Prerequisites: Complete current module plus Flask basics.
- 2. Geospatial Visualization** Choropleth maps, interactive zoom/pan. Dataset: COVID-19 or climate data. New concepts: Projections, spatial aggregation.
- 3. Network Visualization** Graph layouts (force-directed, hierarchical). Tools: NetworkX plus Plotly. Use case: Social networks, org charts.
- 4. Advanced Storytelling** Scrollytelling (scrolling narratives). Animation sequencing. Tool: Reveal.js integration.

4.4.3 Long-Term Vision

- 1. Adaptive Learning Paths** Pre-assessment quiz determines starting point. Beginner track (more scaffolding) vs. advanced track (more independence). Powered by learning analytics and ML recommendation.
- 2. Domain-Specific Modules** Healthcare: Clinical dashboards, patient outcomes Finance: Time series, portfolio visualization Marketing: Funnel analysis, cohort retention Each uses same principles, different datasets.
- 3. Capstone Integration** Final project spans multiple modules. Requirements: ETL plus analysis plus visualization plus deployment. Simulates end-to-end data science workflow.
- 4. Industry Partnerships** Guest speakers from visualization roles (Bloomberg, NYT graphics). Real client projects (anonymized datasets). Internship pathways.

4.4.4 Assessment of Current Effectiveness

Metrics for Evaluation:

Concept Mastery:

- Target: 80% score Bloom's "Apply" or higher
- Measurement: Exercise 2-3 performance

Engagement:

- Target: >90% complete all exercises
- Measurement: Notebook execution logs

Transfer:

- Target: Students use principles in other courses
- Measurement: Survey alumni 6 months later

Satisfaction:

- Target: >4.5/5 average rating
- Measurement: End-of-module survey

Skill Retention:

- Target: >70% recall after 3 months
- Measurement: Follow-up quiz in capstone

Continuous Improvement Process:

Step 1: Collect Data (surveys, logs, performance) Step 2: Analyze Patterns (common errors, bottlenecks) Step 3: Hypothesize Improvements (add videos, simplify Exercise 2) Step 4: A/B Test (treatment vs. control groups) Step 5: Measure Impact (did change improve outcomes?) Step 6: Iterate (repeat cycle)

Planned A/B Tests:

Hypothesis 1: Adding video walkthroughs reduces confusion on log scales

- Treatment: Video plus text explanation
- Control: Text only
- Measure: Post-quiz score, time-to-complete Exercise 2

Hypothesis 2: Peer review improves Exercise 3 quality

- Treatment: Students review 2 peers before final submission
- Control: No peer review
- Measure: Rubric scores, revision depth

5. CONCLUSION

5.1 Pedagogical Philosophy Summary

This module embodies a perception-first, practice-driven approach to data visualization education. By grounding design decisions in cognitive science rather than aesthetics, students develop transferable principles that endure beyond any specific tool.

Core Tenets:

1. Evidence-Based Design: Every choice justified by perceptual research
2. Progressive Complexity: Scaffolding from simple to sophisticated
3. Authentic Assessment: Storytelling mirrors real-world communication
4. Multimodal Learning: Visual, kinesthetic, and textual approaches

5.2 Measured Impact

Expected Outcomes:

- Students create visualizations that are more accurate (correct chart types)
- Students create visualizations that are more accessible (colorblind-friendly)
- Students articulate design rationales (not just "looks good")
- Students integrate storytelling with technical analysis

5.3 Contribution to Field

This module advances data visualization pedagogy by:

1. Explicitly teaching perceptual science (often omitted in technical courses)
2. Balancing theory and practice (many courses focus only on tools)
3. Emphasizing storytelling as core skill (not afterthought)
4. Providing reproducible, open-source teaching materials

5.4 Final Reflection

Effective data visualization is both science (perceptual principles) and art (communication craft). This module demonstrates that both can be taught systematically through evidence-based pedagogy, scaffolded practice, and authentic assessment.

The ultimate goal: Empower students to transform data into insights, and insights into action.

REFERENCES

1. Cleveland, W. S., & McGill, R. (1984). Graphical perception: Theory, experimentation, and application to the development of graphical methods. *Journal of the American Statistical Association*, 79(387), 531-554.
2. Tufte, E. R. (1983). *The visual display of quantitative information*. Graphics Press.
3. Ware, C. (2012). *Information visualization: Perception for design* (3rd ed.). Morgan Kaufmann.
4. Hullman, J., Adar, E., & Shah, P. (2011). Benefitting InfoVis with visual difficulties. *IEEE Transactions on Visualization and Computer Graphics*, 17(12), 2213-2222.
5. Wiggins, G., & McTighe, J. (2005). *Understanding by design* (2nd ed.). ASCD.
6. Few, S. (2012). *Show me the numbers: Designing tables and graphs to enlighten* (2nd ed.). Analytics Press.
7. Knaflic, C. N. (2015). *Storytelling with data: A data visualization guide for business professionals*. Wiley.
8. Munzner, T. (2014). *Visualization analysis and design*. CRC Press.

DOCUMENT METADATA

Author: Mubin Modi

Course: INFO 7390 – Advanced Data Science and Architecture

Module: Advanced Data Visualization

Version: 1.0

Date: December 2025

Page Count: 10 pages

Word Count: Approximately 7,200 words

END OF PEDAGOGICAL REPORT